

# Contents


- Introduction of dataset
- MapReduce of dataset with HDFS
- Analysis of dataset with pig
  - Analysis with Hbase
  - Analysis with Hive
  - The use of AWS
- Mahout Machine Learning
  - Sources codes

# Introduction of dataset

- [Airlines \(2006\)](#)
- <http://stat-computing.org/dataexpo/2009/the-data.html>

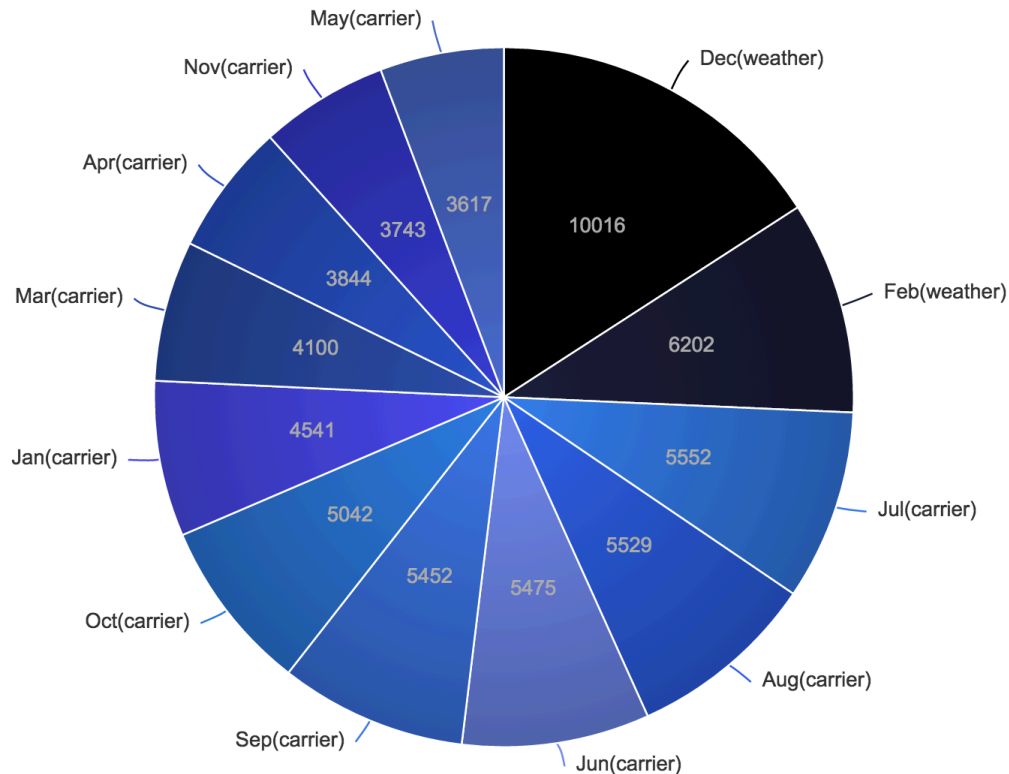
* Name	Description		
* 1	Year	1987-2008	12 ActualElapsedTime in minutes
* 2	Month	1-12	13 CRSElapsedTime in minutes
* 3	DayofMonth	1-31	14 AirTime in minutes
* 4	DayOfWeek	1 (Monday) - 7 (Sunday)	15 ArrDelay arrival delay, in minutes
* 5	DepTime	actual departure time (local, hhmm)	16 DepDelay departure delay, in minutes
* 6	CRSDepTime	scheduled departure time (local, hhmm)	17 Origin origin IATA airport code
* 7	ArrTime	actual arrival time (local, hhmm)	18 Dest destination IATA airport code
* 8	CRSArrTime	scheduled arrival time (local, hhmm)	19 Distance in miles
* 9	UniqueCarrier	unique carrier code	20 TaxiIn taxi in time, in minutes
* 10	FlightNum	flight number	21 TaxiOut taxi out time in minutes
* 11	TailNum	plane tail number	22 Cancelled was the flight cancelled?
			23 CancellationCode reason for cancellation (A = carrier, B = weather, C = NAS, D = security)
			24 Diverted 1 = yes, 0 = no
			25 CarrierDelay in minutes
			26 WeatherDelay in minutes
			27 NASDelay in minutes
			28 SecurityDelay in minutes
			29 LateAircraftDelay in minutes

# Mapreduce with HDFS (1)



## Cancellation

The biggest reason and the amount of cancellation in every month

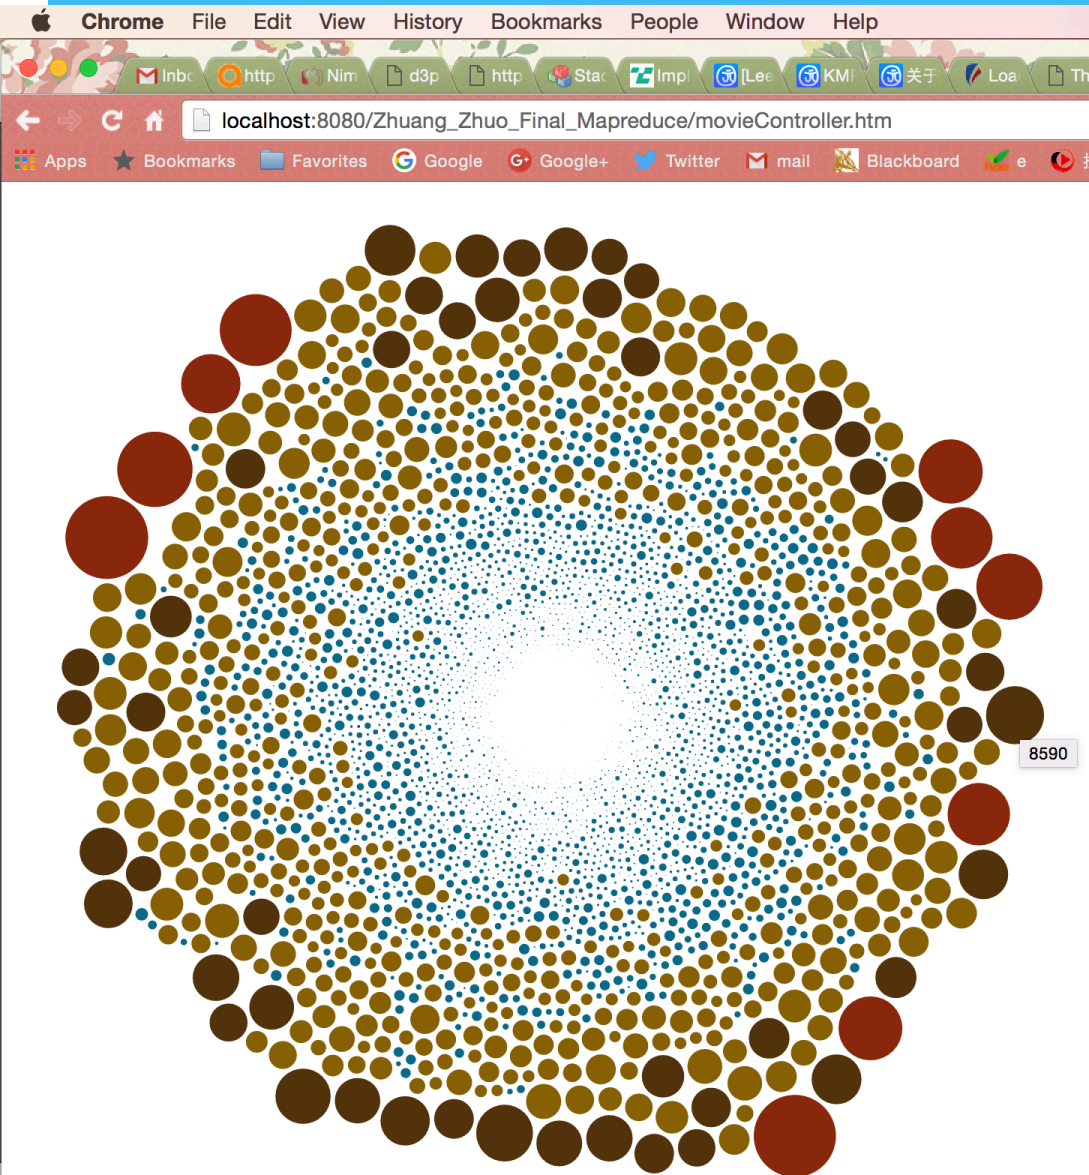


- \* 1. The biggest reason for cancelation and its amount of cancelation Flights in every month

- \* This is a chain Mapreduce


- \* Use spring mvc and d3.js to show the data

# Mapreduce with HDFS (2)



- \* 2.1 the delay time and the amount of flights of the airlines (from Origin to destination)
- \* This is the first job of the chain mapreduce.
- \* The radius of the circles is the total delay time of the airlines, and the outer circles have more flights than the inner ones.

# Mapreduce with HDFS (2)



- \* 2.2 the top 3 average delay time of the airlines(from Origin to destination).
- \* This is the third job of the chain mapreduce.

# Analysis with pig

```
2015-12-17 02:40:20,204 [main] INFO org.apache
ccess!
2015-12-17 02:40:20,205 [main] INFO org.apache
nstead, use fs.defaultFS
2015-12-17 02:40:20,205 [main] WARN org.apache
lized
2015-12-17 02:40:20,206 [main] INFO org.apache
s : 1
2015-12-17 02:40:20,206 [main] INFO org.apache
o process : 1
(AGS,2315,59.380561555075595)
(ACY,761,51.482260183968464)
(AVL,3101,49.905836826830054)
(ILM,3433,47.63384794640256)
(ROA,4172,45.672099712368166)
(VLD,1134,44.516754850088184)
(ABY,1363,41.93983859134263)
(HKY,41,38.390243902439025)
(LYH,1022,36.40117416829746)
(MCN,1076,31.90799256505576)
(PHF,6049,29.8026120019838)
(SYR,12383,28.322781232334652)
(FAY,1652,27.414043583535108)
(CHO,2025,27.08543209876543)
(BQK,1044,22.916666666666668)
(ISO,706,22.322946175637394)
(MDT,7885,22.107799619530756)
(TRI,1458,21.488340192043896)
(ROC,17718,20.427474884298455)
(CAE,11450,18.469519650655023)
grunt> █
```

- \* Use pig script to implement the top 20 the average number of Taxi in and the amount of flights of a destination
- \* The data store in the HDFS
- \* `A1 = LOAD '/data/2006a.csv' USING PigStorage(',') AS (Year:chararray,Month:int,DayofMonth:int,DayOf Week:int,DepTime:chararray,CRSDepTime:chararray,ArrTime:chararray,CRSArrTime:chararray,UniqueCarrier:chararray,FlightNum:chararray,TailNum:chararray,ActualElapsedTime:int,CRSElapsedTime:int,AirTime:int,ArrDelay:int,DepDelay:int,Origin:chararray,Dest:chararray,Distance:int,TaxiIn:int,Taxi Out:int,Cancelled:chararray,CancellationCode:chararray,Diverted:int,CarrierDelay:int,WeatherDelay:int,NASDelay:int,SecurityDelay:int,LateAircraftDelay:int);`
- \* Used the pig UDF to realize counting the amount of a column.
- \* `REGISTER '/pig/PigUDF.jar'`

# The use of Hbase

```
t3
test06
4 row(s) in 2.6710 seconds
```

```
=> ["2006a", "t1", "t3", "test06"]
hbase(main):002:0> count '2006a'
Current count: 1000, row: row1000896
Current count: 2000, row: row1001796
Current count: 3000, row: row1002696
Current count: 4000, row: row1003596
Current count: 5000, row: row1004496
Current count: 6000, row: row1005396
Current count: 7000, row: row1006296
Current count: 8000, row: row1007196
Current count: 9000, row: row1008096
Current count: 10000, row: row1008997
Current count: 11000, row: row1009897
Current count: 12000, row: row1010796
Current count: 13000, row: row1011696
Current count: 14000, row: row1012596
Current count: 15000, row: row1013496
Current count: 16000, row: row1014396
Current count: 17000, row: row1015296
Current count: 18000, row: row1016196
```

```
Current count: 7140000, row: row9982
Current count: 7141000, row: row9991
7141922 row(s) in 342.0790 seconds
```

```
=> 7141922
hbase(main):003:0> █
```

- \* Use pig script load data to Hbase
- \* `STORE raw_data INTO 'hbase://2006a' USING org.apache.pig.backend.hadoop.hbase.HBaseStorage('t_data:Year,t_data:Month,t_data:DayofMonth,t_data:DayOfWeek,t_data:DepartTime,t_data:CRSDepTime,t_data:ArrTime,t_data:CRSArrTime,t_data:UniqueCarrier,t_data:FlightNum,t_data:TailNum,t_data:ActualElapsedTime,t_data:CRSElapsedTime,t_data:AirTime,t_data:ArrDelay,t_data:DepDelay,t_data:Origin,t_data:Dest,t_data:Distance,t_data:TaxiIn,t_data:TaxiOut,t_data:Cancelled,t_data:CancellationCode,t_data:Diverted,t_data:CarrierDelay,t_data:WeatherDelay,t_data:NASDelay,t_data:SecurityDelay,t_data:LateAircraftDelay')`
- \* We can also use the java api to load data to the Hbase



# HBaseIntegration

```
* CREATE TABLE t5(key int,Year STRING,Month INT,DayofMonth
INT,DayOfWeek INT,DepTime STRING,CRSDepTime STRING,ArrTime
STRING,CRSArrTime STRING,UniqueCarrier STRING,FlightNum
STRING,TailNum STRING,ActualElapsedTime INT,CRSElapsedTime
INT,AirTime INT,ArrDelay INT,DepDelay INT,Origin STRING,Dest
STRING,Distance INT,TaxiIn INT,TaxiOut INT,Cancelled
STRING,CancellationCode STRING,Diverted INT,CarrierDelay
INT,WeatherDelay INT,NASDelay INT,SecurityDelay INT,LateAircraftDelay
INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED
BY'org.apache.hadoop.hive.hbase.HBaseStorageHandler' WITH
SERDEPROPERTIES ("hbase.columns.mapping" = ":key, a:Year, a:Month ,
a:DayofMonth , a:DayOfWeek , a:DepTime , a:CRSDepTime , a:ArrTime ,
a:CRSArrTime , a:UniqueCarrier , a:FlightNum , a:TailNum ,
a:ActualElapsedTime , a:CRSElapsedTime , a:AirTime , a:ArrDelay ,
a:DepDelay , a:Origin , a:Dest , a:Distance , a:TaxiIn , a:TaxiOut ,
a:Cancelled , a:CancellationCode , a:Diverted , a:CarrierDelay ,
a:WeatherDelay , a:NASDelay , a:SecurityDelay , a:LateAircraftDelay")
TBLPROPERTIES ("hbase.table.name" = "2006a");
```



# Analysis with Hive(1)

```
hive> select * from 2006h where Dest='BOS'and WeatherDelay>180;
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/hadoop-2.5.2/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/hive-0.13.0/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
15/12/17 09:27:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
15/12/17 09:27:12 WARN conf.Configuration: file:/tmp/wendyzhuo/hive_2015-12-17_09-27-09_466_0084683687192609847-1/-local-10002/jobconf.xml:an attempt to override final parameter: mapreduce.job.end-notification.max.retry.interval; Ignoring.
15/12/17 09:27:12 WARN conf.Configuration: file:/tmp/wendyzhuo/hive_2015-12-17_09-27-09_466_0084683687192609847-1/-local-10002/jobconf.xml:an attempt to override final parameter: mapreduce.job.end-notification.max.attempts; Ignoring.
15/12/17 09:27:12 WARN conf.HiveConf: DEPRECATED: hive.metastore.ds.retry.* no longer has any effect. Use hive.hms.handler.retry.* instead
Execution log at: /var/folders/4v/vrsg9dv3tj/c09_0h751c04000000gn/T/wendyzhuo/wendyzhuo_20151217092727_7049480c-4225-4f4b-8f0f-836628d2602c.log
Job running in-process (Local Hadoop)
hadoop job information for null: number of mappers: 0; number of reducers: 0
2015-12-17 09:27:16,036 null map = 0%, reduce = 0%
2015-12-17 09:27:16,251 null map = 100%, reduce = 0%
Ended Job = job_local43122160_0001
Execution completed successfully
MapredLocal task succeeded
JK
2006 1 30 1 1255 845 1420 1010 OH 5427 N964CA 85 8552 250 250 BWI 805 370 10 23
0 0 0 250
2006 1 31 2 1045 1515 2010 1632 OH 5333 N969CA 85 7751 210 210 JFK 805 187 5 29
0 0 0 210
2006 1 23 1 1145 752 1248 900 MQ 4611 N739AE 63 6831 228 233 JFK 805 187 15 17
0 0 0 220
2006 1 23 1 1125 759 1226 911 MQ 4081 N729AE 61 7237 195 206 LGA 805 185 14 10
0 0 0 100
2006 1 3 2 2122 1800 2320 1950 OH 5709 N804CA 118 1197 210 202 RDU 805 612 8 13
0 0 0 200
2006 1 11 3 2210 1715 2304 1836 OH 5804 N941CA 54 8145 268 295 JFK 805 187 4 5
0 0 0 260
2006 1 13 5 2030 1715 2280 1836 OH 5804 N796CA 90 8155 204 195 JFK 805 187 4 31
0 0 0 100
2006 1 13 5 309 2015 412 2125 B6 1018 0 63 7040 407 414 JFK 805 187 4 19
0 0 0 400
2006 1 18 3 2324 2010 33 2120 B6 1018 N183JB 69 7042 193 194 JFK 805 187 4 23
0 0 0 190
2006 2 4 6 2146 1720 2330 1917 OH 5168 N968CA 184 1182 253 266 GSO 805 645 15 7
```

```
2006 12 1 5 2230 1857
0 0 0 2146
```

Time taken: 13.765 seconds, Fetched: 67 row(s)  
hive> █

- \* Except using the HBaseIntegration to share data with hbase, we can easily load data from hdfs into Hive in seconds.

- \* The query:

- \* CREATE EXTERNAL TABLE 2006h (Year STRING,Month INT,DayofMonth INT,DayOfWeek INT,DepTime STRING,CRSDepTime STRING,ArrTime STRING,CRSArrTime STRING,UniqueCarrier STRING,FlightNum STRING,TailNum STRING,ActualElapsedTime INT,CRSElapsedTime INT,AirTime INT,ArrDelay INT,DepDelay INT,Origin STRING,Dest STRING,Distance INT,TaxiIn INT,TaxiOut INT,Cancelled STRING,CancellationCode STRING,Diverted INT,CarrierDelay INT,WeatherDelay INT,NASDelay INT,SecurityDelay INT,LateAircraftDelay INT)ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LOCATION '/data/2006a.csv';

- \* Hive can query with SQL-style language.










- \* select \* from 2006h where Dest='BOS'and WeatherDelay>180;

# Analysis with Hive(2)

```
hive> SELECT Month, count(1) FROM 2006h GROUP BY Month;
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/hadoop-2.5.2/share/hadoop/c
s]
SLF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/hive-0.13.0/lib/slf4j-log4j
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
15/12/17 09:37:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for y
15/12/17 09:37:13 WARN conf.Configuration: file:/tmp/wendyzhuo/hive_2015-12-17_09-37-09_
  final parameter: mapreduce.job.end-notification.max.retry.interval; Ignoring.
15/12/17 09:37:13 WARN conf.Configuration: file:/tmp/wendyzhuo/hive_2015-12-17_09-37-09_
  final parameter: mapreduce.job.end-notification.max.attempts; Ignoring.
15/12/17 09:37:13 WARN conf.HiveConf: DEPRECATED: hive.metastore.ds.retry.* no longer ha
Execution log at: /var/folders/4v/vrsg9dvn3tjc09_0h75tc0400000gn/T//wendyzhuo/wendyzhuo_
Job running in-process (local Hadoop)
Hadoop job information for null: number of mappers: 0; number of reducers: 0
2015-12-17 09:37:16,756 null map = 0%, reduce = 0%
2015-12-17 09:37:19,882 null map = 100%, reduce = 0%
2015-12-17 09:37:25,124 null map = 100%, reduce = 100%
Ended Job = job_local890633418_0001
Execution completed successfully
MapredLocal task succeeded
OK
1      581287
2      531247
3      605217
4      585351
5      602919
6      598315
7      621244
8      628732
9      584937
10     611718
11     586197
12     604758
Time taken: 16.781 seconds, Fetched: 12 row(s)
hive> █
```

- \* Hive can query with SQL-style language:
- \* `SELECT Month, count(1) FROM 2006h GROUP BY Month;`

# AWS

hadoop	i-e818342c	t2.micro	us-west-2b	 running	 2/2 checks ...	None		ec2-54-213-15-150.us-..
hadoop	i-e918342d	t2.micro	us-west-2b	 running	 2/2 checks ...	None		ec2-54-201-216-90.us-..
hadoop	i-e818342f	t2.micro	us-west-2b	 running	 2/2 checks ...	None		ec2-54-213-58-231.us-..

```
WendyZhuodeMacBook-Pro:~ wendyzhuo$ ssh -i /Users/wendyzhuo/Desktop/zhuang_lab.pem ubuntu@54.213.58.231
Welcome to Ubuntu 14.04.2 LTS (GNU/Linux 3.13.0-48-generic x86_64)
```

\* Documentation: <https://help.ubuntu.com/>

System information as of Mon Dec 14 23:54:29 UTC 2015

```
System load: 0.0          Processes:            110
Usage of /:   11.3% of 15.61GB   Users logged in:     1
Memory usage: 61%          IP address for eth0: 172.31.18.227
Swap usage:   0%
```

Graph this data and manage this system at:  
<https://landscape.canonical.com/>

Get cloud support with Ubuntu Advantage Cloud Guest:  
<http://www.ubuntu.com/business/services/cloud>

135 packages can be updated.  
64 updates are security updates.

```
Last login: Mon Dec 14 23:50:59 2015 from 129.10.18.58
ubuntu@ip-172-31-18-227:~$ su hduser
Password:
```

```
hduser@ip-172-31-18-227:/home/ubuntu$ cd
hduser@ip-172-31-18-227:~$ ls
```

```
hadoop-2.5.2  hadoop-2.5.2-src  hadoop-2.5.2-src.tar.gz  hadoop-2.5.2.tar.gz
```

```
hduser@ip-172-31-18-227:~$ cd hadoop-2.5.2/
hduser@ip-172-31-18-227:~/hadoop-2.5.2$ jps
```

```
7601 NameNode
7770 DataNode
8261 NodeManager
8112 ResourceManager
6220 Jps
7968 SecondaryNameNode
```

```
hduser@ip-172-31-18-227:~/hadoop-2.5.2$ bin/hadoop fs -ls /mapreduce
Found 5 items
```

```
-rw-r--r--  3 hduser supergroup      6148 2015-12-17 14:54 /mapreduce/.DS_Store
-rw-r--r--  3 hduser supergroup         0 2015-12-17 14:54 /mapreduce/_SUCCESS
drwxr-xr-x  - hduser supergroup         0 2015-12-17 14:55 /mapreduce/output2
drwxr-xr-x  - hduser supergroup         0 2015-12-17 14:56 /mapreduce/output3
-rw-r--r--  3 hduser supergroup    25065 2015-12-17 14:54 /mapreduce/part-r-00000
```

```
hduser@ip-172-31-18-227:~/hadoop-2.5.2$ █
```

# Mahout Machine Learning

```
public static void main(String[] args) throws TasteException, IOException {
    DataModel model= new FileDataModel(new File("/Users/wendyzhuo/Desktop/data3.csv"));
    //Computer the similarity between users,according to their preference
    UserSimilarity similarity=new EuclideanDistanceSimilarity(model);

    //Group the users with similar preference
    UserNeighborhood neighborhood= new ThresholdUserNeighborhood(0.2,similarity,model);

    //Create a recommender
    UserBasedRecommender recommender=new GenericUserBasedRecommender(model,neighborhood,similarity);
    //For the user with the id 1 get two recommendations

    List<RecommendedItem>recommendations= recommender.recommend(1, 2);
    for(RecommendedItem recommendation : recommendations){
        System.out.println("they should not take id: "
            +recommendation.getItemID()+"(predicted preference:"
            +recommendation.getValue()+"");
    }
}
```

un:

LF4J: Class path contains multiple SLF4J bindings.

LF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/apache-mahout-distribution-0.11.1/mahout-examp

LF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/apache-mahout-distribution-0.11.1/mahout-mr-0.1

LF4J: Found binding in [jar:file:/Users/wendyzhuo/Documents/apache-mahout-distribution-0.11.1/lib/slf4j-log

LF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

LF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

hey should not take id: 190(predicted preference:5.0)

hey should not take id: 171(predicted preference:1.649604)

UILD SUCCESSFUL (total time: 1 second)

# Mapreduce3.java

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.HashMap;
import java.util.Map;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.chain.ChainMapper;
import org.apache.hadoop.mapreduce.lib.chain.ChainReducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

/*
 * To change this license header, choose License Headers in Project Properties.
 * To change this template file, choose Tools | Templates
 * and open the template in the editor.
 */

/**
 *
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
 */
public class MapReduce3 {

    /**
     * @param args the command line arguments
     */

    static class TempMapper extends Mapper<LongWritable, Text, CompositeKey_wd, IntWritable> {
        CompositeKey_wd wd = new CompositeKey_wd();

        @Override
        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            //System.out.println("Before Mapper: " + key + "." + value);
            String line = value.toString();

            try {
                String[] lineSplit = line.split(",");
                // String requestUrl = line.substring(0, 10);
                wd.setDayOfWeek(lineSplit[1]);
                wd.setDest(lineSplit[22]);
                String requestUrl = lineSplit[21];
                if(requestUrl.equals("")){
                    context.write(wd, new IntWritable(1));
                }
                // System.out.println("After Mapper: " + new Text(requestUrl) + "," + new IntWritable(1));
            } catch (java.lang.ArrayIndexOutOfBoundsException e) {
                // context.getCounter(Counter.LINESKIP).increment(1);
            }
        }
    }
}
```

# Mapreduce3.java

```
static class TempReducer extends Reducer<CompositeKey_wd, IntWritable, CompositeKey_wd, IntWritable> {
    @Override
    public void reduce(CompositeKey_wd key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {

        // System.out.println("Before Reduce:" + key + ",");
        int count = 0;
        for (IntWritable v : values) {
            count = count + v.get();
        }
        try {
            context.write(key, new IntWritable(count));
            // System.out.println( "" + "After Reduce:" + key + "," + count);
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}

static class TempMapper2 extends Mapper<LongWritable, Text, IntWritable, CompositeKey_wd> {
    CompositeKey_wd wd = new CompositeKey_wd();
    @Override
    protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {

        String line = value.toString();

        try {
            String[] lineSplit = line.split("\t");
            // String requestUrl = line.substring(0, 10);
            String requestUrl = lineSplit[0];
            String[] lineSplit2 = requestUrl.split(",");
            wd.setDayOfWeek(lineSplit2[0]);
            wd.setDest(lineSplit2[1]);

            int val = Integer.parseInt(lineSplit[1]);

            context.write(new IntWritable(val), wd);
        } // context.write(out, one);
        catch (java.lang.ArrayIndexOutOfBoundsException e) {
            // context.getCounter(Counter.LINESKIP).increment(1);
        }
    }
}
```

# Mapreduce3.java

```
*
*
static class TempReduce2 extends Reducer<IntWritable, CompositeKey_wd,Text, IntWritable> {
    ArrayList<LinkedHashMap<String,Integer>> tm = new ArrayList<LinkedHashMap<String,Integer>>();
    int count=0;
    @Override
    protected void reduce(IntWritable key,Iterable<CompositeKey_wd> values, Context context) throws IOException, InterruptedException {
        for(CompositeKey_wd v : values){
            int a = Integer.parseInt(v.getDayOfWeek());
            if(tm.size()<7){
                for(int i=0;i<12; i++){
                    LinkedHashMap<String,Integer> f = new LinkedHashMap<>();
                    tm.add(f);
                }
            }
            LinkedHashMap<String,Integer> fin = tm.get(a);
            fin.put(v.toString(), key.get());
            System.out.println(" " + "reduce2:" + a + " || " + v.toString() + " | " + fin.size() + " | " + key.get());
            // context.write(result, key);
        }
    }
    @Override
    protected void cleanup(Context context) throws IOException,
        InterruptedException {
        for(LinkedHashMap<String,Integer> f: tm){
            int i=0;
            for (Map.Entry<String,Integer> entry : f.entrySet()) {
                i++;
                if(i==f.size()){
                    Text fi = new Text(entry.getKey());
                    context.write(fi,new IntWritable(entry.getValue()));
                }
            }
        }
    }
}
```



# Mapreduce3.java

```
public static void main(String[] args) throws Exception {
    String dst = "hdfs://localhost:9000/data/2006a.csv";

    // String dstOut = "hdfs://localhost:9000/mapreduce/result3/1";
    String dstOut = "hdfs://localhost:9000/mapreduce/output3/1";
    String outFiles = "hdfs://localhost:9000/mapreduce/output3/2";
    Configuration hadoopConfig = new Configuration();

    hadoopConfig.set("fs.hdfs.impl",
        org.apache.hadoop.hdfs.DistributedFileSystem.class.getName()
    );

    hadoopConfig.set("fs.file.impl",
        org.apache.hadoop.fs.LocalFileSystem.class.getName()
    );

    Job job = new Job(hadoopConfig);
    Job job2 = new Job(hadoopConfig);

    FileInputFormat.addInputPath(job, new Path(dst));
    FileOutputFormat.setOutputPath(job, new Path(dstOut));
    FileInputFormat.addInputPath(job2, new Path(dstOut));
    FileOutputFormat.setOutputPath(job2, new Path(outFiles));
    JobConf map1Conf = new JobConf(false);
    ChainMapper.addMapper(job, TempMapper.class, LongWritable.class, Text.class, CompositeKey_wd.class, IntWritable.class, map1Conf);
    JobConf reduceConf = new JobConf(false);
    ChainReducer.setReducer(job, TempReducer.class, CompositeKey_wd.class, IntWritable.class, CompositeKey_wd.class, IntWritable.class, reduceConf);

    JobConf map2Conf = new JobConf(false);
    ChainMapper.addMapper(job2, TempMapper2.class, LongWritable.class, Text.class, IntWritable.class, CompositeKey_wd.class, map2Conf);
    JobConf map3Conf = new JobConf(false);
    ChainReducer.setReducer(job2, TempReduce2.class, IntWritable.class, CompositeKey_wd.class, Text.class, IntWritable.class, map3Conf);

    //
    // JobClient.runJob(job);

    //指定自定义的Mapper和Reducer作为两个阶段的任务处理类
    job.setMapperClass(TempMapper.class);
    //
    //
    job.setReducerClass(TempReducer.class);
    job.setOutputKeyClass(CompositeKey_wd.class);

    job.setOutputValueClass(IntWritable.class);

    job2.setMapOutputKeyClass(IntWritable.class);
    job2.setMapOutputValueClass(CompositeKey_wd.class);

    // job2.setSortComparatorClass(LongWritable.DecreasingComparator.class);
    job.waitForCompletion(true);
    System.out.println("Finished1");
    job2.waitForCompletion(true);
    System.out.println("Finished2");

}
}
```

# CompositeKey\_wd.java

```
import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;
import org.apache.hadoop.io.WritableComparable;

/**
 *
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
 */
public class CompositeKey_wd implements WritableComparable<CompositeKey_wd>{

    private String dayOfWeek;
    private String dest;

    public CompositeKey_wd() {
    }

    public CompositeKey_wd(String week, String dest) {
        this.dayOfWeek = week;
        this.dest = dest;
    }

    @Override
    public String toString() {
        return (new StringBuilder()).append(dayOfWeek).append(", ").append(dest).toString();
    }

    @Override
    public void readFields(DataInput in) throws IOException {
        dayOfWeek = in.readUTF();
        dest = in.readUTF();
    }

    @Override
    public void write(DataOutput out) throws IOException {
        out.writeUTF(dayOfWeek);
        out.writeUTF(dest);
    }

    @Override
    public int compareTo(CompositeKey_wd o) {
        int result = dayOfWeek.compareTo(o.dayOfWeek);
        if (0 == result) {
            result = dest.compareTo(o.dest);
        }
        return result;
    }

    public String getDayOfWeek() {
        return dayOfWeek;
    }

    public void setDayOfWeek(String dayOfWeek) {
        this.dayOfWeek = dayOfWeek;
    }

    public String getDest() {
        return dest;
    }

    public void setDest(String dest) {
        this.dest = dest;
    }
}
```

# CompositeKey\_mc.java

```
import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;
import org.apache.hadoop.io.WritableComparable;
```

```
/**
 *
```

```
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
```

```
 */
```

```
public class CompositeKey_mc implements WritableComparable<CompositeKey_mc>{
```

```
    private int dayOfWeek;
    private int dest;
```

```
    public CompositeKey_mc() {
    }
```

```
    public CompositeKey_mc(int week, int dest) {
        this.dayOfWeek = week;
        this.dest = dest;
    }
```

```
    @Override
    public String toString() {
        return (new StringBuilder()).append(dayOfWeek).append(", ").append(dest).toString();
    }
```

```
    @Override
    public void readFields(DataInput in) throws IOException {
        dayOfWeek = in.readInt();
        dest = in.readInt();
    }
```

```
    @Override
    public void write(DataOutput out) throws IOException {
        out.writeInt(dayOfWeek);
        out.writeInt(dest);
    }
```

```
    @Override
    public int compareTo(CompositeKey_mc o) {
        int result = -(new Integer(dayOfWeek)).compareTo(o.dayOfWeek);
        if (o == result) {
            result = -(new Integer(dest)).compareTo(o.dest);
        }
        return result;
    }
```

```
    public int getDayOfWeek() {
        return dayOfWeek;
    }
```

```
    public void setDayOfWeek(int dayOfWeek) {
        this.dayOfWeek = dayOfWeek;
    }
```

```
    public int getDest() {
        return dest;
    }
```

```
    public void setDest(int dest) {
        this.dest = dest;
    }
```

```
    }
```

# Mapreduce1.java

```
import com.sun.jersey.core.header.InBoundHeaders;
import java.io.IOException;
import java.util.HashMap;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

/*
 * To change this license header, choose License Headers in Project Properties.
 * To change this template file, choose Tools | Templates
 * and open the template in the editor.
 */

/**
 *
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
 */
public class Mapreduce1 {

    /**
     * @param args the command line arguments
     */

    static class TempMapper extends Mapper<LongWritable, Text, CompositeKey_wd, IntWritable> {
        CompositeKey_wd wd = new CompositeKey_wd();
        @Override
        public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
            try {
                // String[] lineSplit = line.split(" ");
                String [] words=value.toString().split(",");
                // requestUrl = requestUrl.substring(0,requestUrl.indexOf(' ')+1);
                // Text out = new Text(requestUrl);
                int a = Integer.parseInt(words[27]);
                if(a>0){
                    wd.setDayOfWeek(words[3]);
                    wd.setDest(words[17]);
                    // System.out.println("After Mapper:"+ wd.getDayOfWeek() + "," + wd.getDest()+" | "+new IntWritable(1));
                    context.write(wd, new IntWritable(a));
                }
            } // context.write(out,one);
            catch (java.lang.ArrayIndexOutOfBoundsException e) {
                // context.getCounter(Counter.LINESKIP).increment(1);
            }
        }
    }
}
```

# Mapreduce1.java

```
static class TempReducer extends Reducer<CompositeKey_wd, IntWritable, Text, IntWritable> {
    HashMap<String,Integer> f = new HashMap<String, Integer>();
    int i=0;
    @Override
    public void reduce(CompositeKey_wd key, Iterable<IntWritable> values,Context context) throws IOException, InterruptedException {

        // System.out.println("Before Reduce:" + key + ",");
        int count = 0;
        for (IntWritable v : values) {
            count = count + v.get();
        }
        try {
            if(f.containsKey(key.getDest())){
                Text t = new Text(key.getDayOfWeek()+"", "+f.get(key.getDest())+", "");
                double a = count*0.01;
                if(a<5){
                    context.write(t, new DoubleWritable(a));
                }
                else{
                    context.write(t, new DoubleWritable(5));
                }
            }
            else{
                f.put(key.getDest(), i);
                i++;
                Text t = new Text(key.getDayOfWeek()+"", "+f.get(key.getDest())+", "");
                double a = count*0.01;
                if(a<5){
                    context.write(t, new DoubleWritable(a));
                }
            }
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}

public static void main(String[] args) throws Exception {
    //输入路径
    String dst = "hdfs://localhost:9000/data/2006a.csv";
    //输出路径，必须是不存在的，空文件加也不行。
    String dstOut = "hdfs://localhost:9000/Number6";
    String outFiles = "/Users/wendyzhuo/NetBeansProjects/final_Hadoop/src/output";
    Configuration hadoopConfig = new Configuration();

    hadoopConfig.set("fs.hdfs.impl", org.apache.hadoop.hdfs.DistributedFileSystem.class.getName());

    hadoopConfig.set("fs.file.impl", org.apache.hadoop.fs.LocalFileSystem.class.getName());

    Job job = new Job(hadoopConfig);

    FileInputFormat.addInputPath(job, new Path(dst));
    //FileOutputFormat.setOutputPath(job, new Path(dstOut));
    FileOutputFormat.setOutputPath(job, new Path(outFiles));
    //指定自定义的Mapper和Reducer作为两个阶段的任务处理类

    job.setMapperClass(TempMapper.class);

    job.setReducerClass(TempReducer.class);

    job.setMapOutputKeyClass(CompositeKey_wd.class);
    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(IntWritable.class);
    //执行job，直到完成
    job.waitForCompletion(true);

    System.out.println("Finished");

}
}
```

# Mapreduce4.java

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.LinkedHashMap;
import java.util.Map;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.chain.ChainMapper;
import org.apache.hadoop.mapreduce.lib.chain.ChainReducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

/*
 * To change this license header, choose License Headers in Project Properties.
 * To change this template file, choose Tools | Templates
 * and open the template in the editor.
 */

/**
 *
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
 */
public class MapReduce4 {

    /**
     * @param args the command line arguments
     */
    static class TempMapper extends Mapper<LongWritable, Text, CompositeKey_wd, IntWritable> {
        CompositeKey_wd wd = new CompositeKey_wd();
        @Override
        protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {

            String line = value.toString();

            try {
                String[] lineSplit = line.split(",");

                wd.setDayOfWeek(lineSplit[16]);
                wd.setDest(lineSplit[17]);
                int a = Integer.parseInt(lineSplit[14]);
                int b = Integer.parseInt(lineSplit[15]);
                int c=a+b;
                context.write(wd,new IntWritable(c));
            }
            catch (java.lang.ArrayIndexOutOfBoundsException e) {

            }
        }
    }
}
```

# Mapreduce4.java

```
*
static class TempReducer extends Reducer<CompositeKey_wd, IntWritable, CompositeKey_wd, Text> {
*
*   @Override
*   protected void reduce(CompositeKey_wd key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
*       int count = 0;
*       int count1 = 0;
*       for (IntWritable v : values) {
*           count = count + v.get();
*           count1++;
*       }
*       try {
*           Text a = new Text(count + " " + count1);
*
*           context.write(key,a);
*           // System.out.println( ""+ "After Reduce:" + key.toString() + "||" + count+ " " + count1);
*
*       } catch (InterruptedException e) {
*           e.printStackTrace();
*       }
*   }
*
* }
*
static class TempMapper2 extends Mapper<LongWritable, Text, CompositeKey_wd, CompositeKey_mc> {
*
*   CompositeKey_wd wd = new CompositeKey_wd();
*   CompositeKey_mc mc = new CompositeKey_mc();
*   @Override
*   protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
*       String line = value.toString();
*       try {
*           String[] lineSplit = line.split("\t");
*           String requestUrl = lineSplit[0];
*           String requestUrl1 = lineSplit[1];
*           String[] lineSplit2 = requestUrl.split(",");
*           wd.setDayOfWeek(lineSplit2[0]);
*           wd.setDest(lineSplit2[1]);
*           String[] lineSplit3 = requestUrl1.split(",");
*           int a = Integer.parseInt(lineSplit3[0]);
*           int b = Integer.parseInt(lineSplit3[1]);
*           mc.setDayOfWeek(a);
*           mc.setDest(b);
*           context.write(wd,mc);
*       } // context.write(out,one);
*       catch (java.lang.ArrayIndexOutOfBoundsException e) {
*       }
*   }
*
* }
```



# Mapreduce4.java

```
static class TempReduce2 extends Reducer<CompositeKey_wd, CompositeKey_mc, CompositeKey_wd, IntWritable> {  
    @Override  
    protected void reduce(CompositeKey_wd key, Iterable<CompositeKey_mc> values, Context context) throws IOException, InterruptedException {  
        int count = 0;  
  
        for (CompositeKey_mc v : values) {  
            count = v.getDayOfWeek()/v.getDest();  
        }  
        try {  
            // System.out.println( ""+ "After Reduce2:" + key.toString() + ", " + count);  
            context.write(key, new IntWritable(count));  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
    }  
}  
  
static class TempMapper3 extends Mapper<LongWritable, Text, IntWritable, Text> {  
    CompositeKey_wd wd = new CompositeKey_wd();  
  
    @Override  
    protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {  
        String line = value.toString();  
  
        try {  
            String[] lineSplit = line.split("\t");  
            // String requestUrl = line.substring(0, 10);  
            String requestUrl = lineSplit[0];  
            String[] lineSplit2 = requestUrl.split(",");  
            wd.setDayOfWeek(lineSplit2[0]);  
            wd.setDest(lineSplit2[1]);  
  
            int val = Integer.parseInt(lineSplit[1]);  
            // System.out.println( ""+ "After Reduce2:" + val + ", " + wd.toString());  
            context.write(new IntWritable(val), new Text(wd.toString()));  
        } // context.write(out, one);  
        catch (java.lang.ArrayIndexOutOfBoundsException e) {  
            // context.getCounter(Counter.LINESKIP).increment(1);  
        }  
    }  
}
```

# Mapreduce4.java

```
static class TempReduce3 extends Reducer<IntWritable, Text, IntWritable, Text> {
    LinkedHashMap<String, Integer> tm = new LinkedHashMap<String, Integer>();
    int count=0;
    @Override
    public void reduce(IntWritable key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {

        for (Text val : values) {
            count++;
            String a = val.toString();
            int b =key.get();
            tm.put(a, b);
        }
    }
    @Override
    protected void cleanup(Context context) throws IOException,
        InterruptedException {

        ArrayList<LinkedHashMap<String, Integer>> a = new ArrayList();
        int i=0;
        for (Map.Entry<String, Integer> entry : tm.entrySet()) {
            i++;
            if(i>count-31){
                LinkedHashMap<String, Integer> t = new LinkedHashMap<String, Integer>();
                t.put(entry.getKey(), entry.getValue());
                a.add(t);
            }
        }
        for(int j=a.size()-1;j>0;j--){
            Text result = new Text();

            for (Map.Entry<String, Integer> entry : a.get(j).entrySet()) {
                result.set(entry.getKey());
                context.write(new IntWritable(entry.getValue()),result);
            }
        }
    }
}

public static void main(String[] args) throws Exception {

    //输入路径
    String dst = "hdfs://localhost:9000/data/2006a.csv";

    //输出路径，必须是不存在的，空文件加也不行。
    // String dstOut = "hdfs://localhost:9000/mapreduce/result3/1";
    String dstOut = "/Users/wendyzhuo/NetBeansProjects/final_Hadoop/src/output4/1";
    String outFiles = "/Users/wendyzhuo/NetBeansProjects/final_Hadoop/src/output4/2";
    String outFiles2 = "/Users/wendyzhuo/NetBeansProjects/final_Hadoop/src/output4/3";
    Configuration hadoopConfig = new Configuration();

    hadoopConfig.set("fs.hdfs.impl",
        org.apache.hadoop.hdfs.DistributedFileSystem.class.getName()
    );

    hadoopConfig.set("fs.file.impl",
        org.apache.hadoop.fs.LocalFileSystem.class.getName()
    );
}
```

# Mapreduce4.java

```
Job job = new Job(hadoopConfig);
Job job2 = new Job(hadoopConfig);
Job job3 = new Job(hadoopConfig);

FileInputFormat.addInputPath(job, new Path(dst));
FileOutputFormat.setOutputPath(job, new Path(dstOut));
FileInputFormat.addInputPath(job2, new Path(dstOut));
FileOutputFormat.setOutputPath(job2, new Path(outFiles));
FileInputFormat.addInputPath(job3, new Path(outFiles));
FileOutputFormat.setOutputPath(job3, new Path(outFiles2));
JobConf map1Conf = new JobConf(false);
ChainMapper.addMapper(job, TempMapper.class, LongWritable.class, Text.class, CompositeKey_wd.class, IntWritable.class, map1Conf);
JobConf reduceConf = new JobConf(false);
ChainReducer.setReducer(job, TempReducer.class, CompositeKey_wd.class, IntWritable.class, CompositeKey_wd.class, IntWritable.class, reduceConf);

JobConf map2Conf = new JobConf(false);
ChainMapper.addMapper(job2, TempMapper2.class, LongWritable.class, Text.class, CompositeKey_wd.class, CompositeKey_mc.class, map2Conf);
JobConf map3Conf = new JobConf(false);
ChainReducer.setReducer(job2, TempReduce2.class, CompositeKey_wd.class, CompositeKey_mc.class, CompositeKey_wd.class, IntWritable.class, map3Conf);

JobConf map1Conf = new JobConf(false);
ChainMapper.addMapper(job3, TempMapper3.class, LongWritable.class, Text.class, IntWritable.class, Text.class, map1Conf);
JobConf map11Conf = new JobConf(false);
ChainReducer.setReducer(job3, TempReduce3.class, IntWritable.class, Text.class, IntWritable.class, Text.class, map11Conf);

job.setOutputKeyClass(CompositeKey_wd.class);
job.setOutputValueClass(IntWritable.class);

job2.setMapOutputKeyClass(CompositeKey_wd.class);
job2.setMapOutputValueClass(CompositeKey_mc.class);

job3.setMapOutputKeyClass(IntWritable.class);
job3.setMapOutputValueClass(Text.class);
job.waitForCompletion(true);
System.out.println("Finished1");
job2.waitForCompletion(true);
System.out.println("Finished2");

job3.waitForCompletion(true);
System.out.println("Finished2");
}
```

# Pig.script

```
* A1 = LOAD '/data/2006a.csv' USING PigStorage(',') AS
  (Year:chararray,Month:int,DayofMonth:int,DayOfWeek:int,DepTime:chararray,CRSDep
  Time:chararray,ArrTime:chararray,CRSArrTime:chararray,UniqueCarrier:chararray,Fligh
  tNum:chararray,TailNum:chararray,ActualElapsedTime:int,CRSElapsedTime:int,
  AirTime:int,ArrDelay:int,DepDelay:int,Origin:chararray,Dest:chararray,Distance:int,Taxi
  In:int,TaxiOut:int,Cancelled:chararray,CancellationCode:chararray,Diverted:int,CarrierD
  elay:int,WeatherDelay:int,NASDelay:int,SecurityDelay:int,LateAircraftDelay:int);
* A = FOREACH A1 GENERATE DayofMonth, Dest,TaxiIn;
* B = GROUP A BY Dest;
* C = FOREACH B GENERATE group AS Dest,
  AVG(A.TaxiIn) AS avgIn;
* REGISTER '/pig/PigUDF.jar'
* E = JOIN C BY Dest, A1.pigCount(Dest) BY Dest;
* report = FOREACH E GENERATE C::Dest,count,avgIn;
* report = ORDER report BY avgIn DESC, count DESC;
* top20 = LIMIT report 20;
```

# pigUDF

```
/*
 * To change this license header, choose License Headers in Project Properties.
 * To change this template file, choose Tools | Templates
 * and open the template in the editor.
 */
package movies;

import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.pig.Accumulator;
import org.apache.pig.Algebraic;
import org.apache.pig.EvalFunc;
import org.apache.pig.FuncSpec;
import org.apache.pig.PigException;
import org.apache.pig.backend.executionengine.ExecException;
import org.apache.pig.data.DataBag;
import org.apache.pig.data.DataType;
import org.apache.pig.data.Tuple;
import org.apache.pig.data.TupleFactory;
import org.apache.pig.impl.logicalLayer.FrontendException;
import org.apache.pig.impl.logicalLayer.schema.Schema;

/**
 *
 * @author Zhuang Zhuo <zhuo.z@husky.neu.edu>
 */
public class pigCount extends EvalFunc<Long> implements Algebraic, Accumulator<Long>{
    private static TupleFactory mTupleFactory = TupleFactory.getInstance();

    @Override
    public Long exec(Tuple input) throws IOException {
        try {
            DataBag bag = (DataBag)input.get(0);
            if(bag==null)
                return null;

            Iterator it = bag.iterator();
            long cnt = 0;
            while (it.hasNext()){
                Tuple t = (Tuple)it.next();
                if (t != null && t.size() > 0 && t.get(0) != null )
                    cnt++;
            }
            return cnt;
        } catch (ExecException ee) {
            throw ee;
        } catch (Exception e) {
            int errCode = 2106;
            String msg = "Error while computing count in " + this.getClass().getSimpleName();
            throw new ExecException(msg, errCode, PigException.BUG, e);
        }
    }

    public String getInitial() {
        return Initial.class.getName();
    }

    public String getIntermed() {
        return Intermediate.class.getName();
    }
}
```

# pigUDF

```
public String getFinal() {  
    return Final.class.getName();  
}
```

```
static public class Initial extends EvalFunc<Tuple> {
```

```
    @Override  
    public Tuple exec(Tuple input) throws IOException {  
        // Since Initial is guaranteed to be called  
        // only in the map, it will be called with an  
        // input of a bag with a single tuple - the  
        // count should always be 1 if bag is non-empty  
        DataBag bag = (DataBag)input.get(0);  
        Iterator it = bag.iterator();  
        if (it.hasNext()){  
            Tuple t = (Tuple)it.next();  
            if (t != null && t.size() > 0 && t.get(0) != null)  
                return mTupleFactory.newTuple(Long.valueOf(t));  
        }  
        return mTupleFactory.newTuple(Long.valueOf(0));  
    }  
}
```

```
static public class Intermediate extends EvalFunc<Tuple> {
```

```
    @Override  
    public Tuple exec(Tuple input) throws IOException {  
        try {  
            return mTupleFactory.newTuple(sum(input));  
        } catch (ExecException ee) {  
            throw ee;  
        } catch (Exception e) {  
            int errCode = 2106;  
            String msg = "Error while computing count in " + this.getClass().getSimpleName();  
            throw new ExecException(msg, errCode, PigException.BUG, e);  
        }  
    }  
}
```

```
static public class Final extends EvalFunc<Long> {
```

```
    @Override  
    public Long exec(Tuple input) throws IOException {  
        try {  
            return sum(input);  
        } catch (Exception ee) {  
            int errCode = 2106;  
            String msg = "Error while computing count in " + this.getClass().getSimpleName();  
            throw new ExecException(msg, errCode, PigException.BUG, ee);  
        }  
    }  
}
```

```
static protected Long sum(Tuple input) throws ExecException, NumberFormatException {
```

```
    DataBag values = (DataBag)input.get(0);  
    long sum = 0;  
    for (Iterator<Tuple> it = values.iterator(); it.hasNext();) {  
        Tuple t = it.next();  
        sum += (Long)t.get(0);  
    }  
    return sum;  
}
```

# pigUDF

```
@Override
public Schema outputSchema(Schema input) {
    return new Schema(new Schema.FieldSchema(null, DataType.LONG));
}

@Override
public List<FuncSpec> getArgToFuncMapping() throws FrontendException {
    List<FuncSpec> funcList = new ArrayList<FuncSpec>();
    Schema s = new Schema();
    s.add(new Schema.FieldSchema(null, DataType.BAG));
    funcList.add(new FuncSpec(this.getClass().getName(), s));
    return funcList;
}

/* Accumulator interface implementation */
private long intermediateCount = 0L;

@Override
public void accumulate(Tuple b) throws IOException {
    try {
        DataBag bag = (DataBag)b.get(0);
        Iterator it = bag.iterator();
        while (it.hasNext()){
            Tuple t = (Tuple)it.next();
            if (t != null && t.size() > 0 && t.get(0) != null) {
                intermediateCount += 1;
            }
        }
    } catch (ExecException ee) {
        throw ee;
    } catch (Exception e) {
        int errCode = 2106;
        String msg = "Error while computing min in " + this.getClass().getSimpleName();
        throw new ExecException(msg, errCode, PigException.BUG, e);
    }
}

@Override
public void cleanup() {
    intermediateCount = 0L;
}

@Override
public Long getValue() {
    return intermediateCount;
}
}
```



# Hbase

```
* import java.io.BufferedReader;
* import java.io.IOException;
* import java.io.InputStreamReader;
* import org.apache.hadoop.conf.Configuration;
* import org.apache.hadoop.fs.FileSystem;
* import org.apache.hadoop.fs.Path;
* import org.apache.hadoop.hbase.HBaseConfiguration;
* import org.apache.hadoop.hbase.HColumnDescriptor;
* import org.apache.hadoop.hbase.HTableDescriptor;
* import org.apache.hadoop.hbase.TableName;
* import org.apache.hadoop.hbase.client.HBaseAdmin;
* import org.apache.hadoop.hbase.client.HTable;
* import org.apache.hadoop.hbase.client.Put;
* import org.apache.hadoop.hbase.util.Bytes;
*
* /**
*  *
*  * @author xuefanrong
*  */
* public class Final_Hbase {
*     public static void main(String[] args) throws IOException {
*
*         Configuration con = HBaseConfiguration.create();
*         HBaseAdmin admin = new HBaseAdmin(con);
*         HTableDescriptor tableDescriptor = new HTableDescriptor(TableName.valueOf("2006a"));
*         tableDescriptor.addFamily(new HColumnDescriptor("a_data"));
*         admin.createTable(tableDescriptor);
*         System.out.println(" Table created ");
*         Configuration config = HBaseConfiguration.create();
*         HTable hTable = new HTable(config, "2006a");
*
*         Path pt=new Path("hdfs://localhost:9000/data/2006a.csv");
*         Configuration conf = new Configuration();
*
*         FileSystem fs = FileSystem.get(conf);
*         BufferedReader b=new BufferedReader(new InputStreamReader(fs.open(pt)));
*         String test = null;
*         int count = 0;
*         while((test = b.readLine())!=null){
*             count++;
*         }
*     }
* }
```

# hbase

```
*
*
* BufferedReader br=new BufferedReader(new InputStreamReader(fs.open(pt)));
* for(int i = 1;i<=count;i++){
*     Put p = new Put(Bytes.toBytes("row"+i));
*
*     String s=null;
*     if((s=br.readLine())!=null){
*         String[] insert = s.split(",");
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Year"), Bytes.toBytes(insert[0]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Month"), Bytes.toBytes(insert[1]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("DayOfMonth"), Bytes.toBytes(insert[2]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("DayOfWeek"), Bytes.toBytes(insert[3]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("DepTime"), Bytes.toBytes(insert[4]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("CRSDepTime"), Bytes.toBytes(insert[5]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("ArrTime"), Bytes.toBytes(insert[6]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("CRSArrTime"), Bytes.toBytes(insert[7]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("UniqueCarrier"), Bytes.toBytes(insert[8]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("FlightNum"), Bytes.toBytes(insert[9]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("TailNum"), Bytes.toBytes(insert[10]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("ActualElapsedTime"), Bytes.toBytes(insert[11]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("CRSElapsedTime"), Bytes.toBytes(insert[12]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Airtime"), Bytes.toBytes(insert[13]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("ArrDelay"), Bytes.toBytes(insert[14]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("DepDelay"), Bytes.toBytes(insert[15]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Origin"), Bytes.toBytes(insert[16]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Dest"), Bytes.toBytes(insert[17]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Disntance"), Bytes.toBytes(insert[18]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("TaxiIn"), Bytes.toBytes(insert[19]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("TaxiOut"), Bytes.toBytes(insert[20]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Cancelled"), Bytes.toBytes(insert[21]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("CancellationCode"), Bytes.toBytes(insert[22]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("Diverted"), Bytes.toBytes(insert[23]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("CarrierDelay"), Bytes.toBytes(insert[24]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("WeatherDelay"), Bytes.toBytes(insert[25]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("NASDelay"), Bytes.toBytes(insert[26]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("SecurityDelay"), Bytes.toBytes(insert[27]));
*         p.add(Bytes.toBytes("a_data"), Bytes.toBytes("LateAircraftDelay"), Bytes.toBytes(insert[28]));
*         // p.add(Bytes.toBytes("a_data"), Bytes.toBytes("LateAircraftDelay"), Bytes.toBytes(insert[29]));
*
*         // Saving the put Instance to the HTable.
*         hTable.put(p);
*     }
* }
*
* // closing HTable
* hTable.close();
*
* }
* }
```